

CS2040C Data Structures and Algorithms

Hashing

For efficient look-up in a table

Outline

- Direct Addressing
- What is hashing?
- Hash Table/Hash Function
- What is collision?
- How to **resolve collision**?
- **Primary** clustering and **secondary** clustering
- STL hash table

Lookup Table

- Most applications require a data structure to:
 - ❑ Store a number of items
 - ❑ Delete a particular item
 - ❑ Search for one particular item using a special piece of information (**key**)
- **Lookup Table** is an abstraction that captures the requirements above
 - ❑ Many different implementations possible!

Lookup Table: Example Implementations

Operations	Unsorted Array / Linked List	Sorted Array	Sorted Linked List
Insert	$O(1)$	$O(N)$	$O(N)$
Delete	$O(N)$	$O(N)$	$O(N)$
Search	$O(N)$	$O(\log_2 N)$	$O(N)$

- Confirm your understanding by verifying the complexity required for each of the operations

Direct Addressing Table

A simplified version of hash table

Example: The SBS Bus Problem

- Consider a system to manage information about **bus services** for the bus company SBS
- The main operations are:
 - We assume bus service number is an integer

Operations	Functionality
<i>Find</i> (N)	Does bus service N exists?
<i>Insert</i> (N)	Add bus service N
<i>Delete</i> (N)	Remove bus service N


Observations: The SBS Bus Problem

- The bus service are indicated by an integer between [1 ... 999]
- Efficient Solution:
 - Use a **boolean array** of 1000 elements
 - Element at index ***N*** represents the bus service ***N***
 - **True** == exists, **False** == not exist
- Known as **direct addressing table**

0	false
1	false
2	true
	:
96	true
	:
151	true
	:
999	false


Direct Addressing Table

- Additional information can be stored
 - E.g. route, interval, number of buses serving for a particular bus service
- Instead of a single boolean value:
 - We can store a **structure/object** at each location



Direct Addressing Table

- Alternatively, we can store a **reference to object in each location**
 - Invalid bus service stores a **NULL**



Generalized Direct Addressing Table

- The generalized set of operations for direct addressing table are:

Operations	Basic Steps	Big-O
<i>Find</i> (N)	return a[N]	O(1)
<i>Insert</i> (N,data)	a[N] = data	O(1)
<i>Delete</i> (N)	a[N] = NULL	O(1)

Direct Addressing Table: Summary

- Direct addressing table is very **efficient**
- However, there are many **restrictions**
 - Key must be **integer** (*what about bus no 95A or NR30?*)
 - Range of keys must be **small**
 - E.g. what if keys are telephone numbers?
 - Keys must be **dense**
 - Most keys in the range are valid
 - Not many "gaps" in the key values

What is Hashing?

- Hashing (system) uses a hash function, a hash table, and a conflict (collision) resolution scheme to implement a table ADT
- A conflict resolution scheme is the action taken to resolve the conflict between two keys that have been assigned the same address by the hash function

Generalized Idea:


- Use a **conversion function** to map:
 - Non-integer to integer
 - Sparse integers in a large range into a dense integers in a smaller range
- This conversion function is known as **hash function**
 - The fundamental idea behind hash table!
 - Hash Table
 - = **Direct Addressing Table** + **Hash Function**

Hash Table

A **generalization** of direct addressing table, to remove its restrictions

Hash table

- Map **large** integers to **smaller** integers
- Map **non-integer** keys to **integers**



Hash Table: Operations

- One additional step:
 - Apply hash function $h()$ to the key value first
 - $h(\text{key})$ gives the **home address** of the key value

Operations	Basic Steps
<i>Find</i> (N)	return $a[h(N)]$
<i>Insert</i> (N,data)	$a[h(N)] = \text{data}$
<i>Delete</i> (N)	$a[h(N)] = \text{NULL}$

- Time complexity now depends on the performance of the **hash function $h()$**

Hash Tables: Problems

- If the result of the hash function is **unique**, each key is mapped to a different home address (array index)
 - known as **perfect hash function**
- This is **not always the case**
 - Given two different keys, it is possible for a hash function to give the same result
 - $\text{Hash}(\text{key1}) == \text{Hash}(\text{key2})$
but $\text{key1} \neq \text{key2}$
- This problem is known as **collision**
 - Need to find ways to resolve them

Hashing Collision

- Given the hash function:

$$h(\text{key}) = \text{key} \% 17$$


$$h(19) = 19 \% 17 = 2$$

$$h(87) = 87 \% 17 = 2$$

Collision!  2



Hash table



A hash function does **not** guarantee that two different keys goes into **different slots**!

E.g. 66753375 hashes to the same location of 66752378.

This is called a "**collision**".

Hash Table: Important Issues

- How to define a good **hash function**?
 - What are the properties of a good hash function?
- How to **resolve collision**?

Hash Functions

Good Hash Function: Properties

- A good hash function should:
 - ❑ Be fast to compute (should be $O(1)$)
 - ❑ Scatter keys evenly throughout the hash table
 - ❑ Result in few or none collisions
 - ❑ Allow the hash table to be small
- These properties should be evaluated in the context of the potential key range

Counter Example

- Selecting digits from several positions usually make poor hashing function
 - E.g. $\text{Hash}(d_0d_1d_2\dots d_7) = d_3d_7$
 - $\text{Hash}(667\mathbf{5}437\mathbf{8}) = \mathbf{58}$
 - $\text{Hash}(634\mathbf{5}909\mathbf{8}) = \mathbf{58}$
- What if we select the first three digits from Singapore phone numbers as the hash value?

Perfect Hash Function

- Perfect Hash Function:
 - One-to-one mapping between the keys and array indices
 - **NO collision**
- It is possible if we know all keys in advance
- **Example:**
 - When a compiler searches for keywords or reserved words

How to Define a Hash Function?

- Uniform hash function
- Division method
- Multiplication method
- Hashing of strings

Uniform Hash Function

- Uniform Hash Function:
 - Distribute the keys **evenly throughout** the hash table
- Formal definition:
 - Given **K keys** and **M locations** in a hash table
 - **$H(K)$** is uniform if each location receives no more than $\left\lceil \frac{K}{M} \right\rceil$ keys

Uniform Hash Function

- Given:

- Keys are integers uniformly distributed in $[0, X)$
- Hash table of size ***m*** ($m < X$)

- We can hash the keys uniformly into the table by:

$$k \in [0, X)$$

$$hash(k) = \left\lfloor \frac{km}{X} \right\rfloor$$

Modulo Method

- Given a hash table of **m slots**
 - We can use the modulo operator "%" to map an integer to a value between **0** and **m-1**:

$$\text{hash}(k) = k \% m$$

- One of the most popular methods
- Behaviour of the hash function depends on:
 - Key distribution
 - Table size **m**

Modulo Method: Table Size m

- Generally, we want the hash function to generate "random-like" home addresses even if the keys are in continuous range
- Some table size should be avoided in modulo method due to commonly encountered key sequence
- **Example:**
 - $m = 10^n$
 - Hash function returns the last n digits of the key!
 - $m = 2^n$
 - Hash function returns the last n bits of the key!

Modulo Method: Table Size m

- Rule of thumb:

- Choose table size to be a **large prime number** close to a power of 2

- Several reasons:

- We can get a "shuffling" effect by first multiplying the key with another prime number q :

$$\text{hash}(k) = (k * q) \% m$$

- Prime table size allows effective collision resolution method (more later)

Multiplicative Method

- Hash function takes the following form:

1. Multiply key with a real number A between $[0..1]$
2. Extract the fractional part
3. Multiply by hash table size, m

$$hash(k) = \lfloor m(kA - \lfloor kA \rfloor) \rfloor$$

- **Rationale:**

- Fraction part of multiplication is "random-like" even for continuous key range

- A common choice for A is the reciprocal of **golden ratio**:

$$A = \frac{\sqrt{5} - 1}{2}$$

Hashing of Strings

- For non-integer keys:
 - We first convert the key into an integer, then apply hash function on the result
- Let us use string as illustration

```
int HashString( string str )
{
    int sum = 0;

    for ( i = 0; i < str.size(); i++ )
        sum += str[i];

    return sum % tableSizeM;
}
```

Convert **str** into an integer **sum**

Perform the actual hashing. Modulo method is used here as example.

Hashing of Strings: Problems

- The method used is not very good:
 - Many strings converted to the **same sum**
 - ➔ Results in large number of collisions
- Example:
 - `HashString("abc") == HashString("bac") == HashString("cba")`
- Problem:
 - The conversion fails to take the **position of each character into account**
 - ➔ Permutation of a string gives the same sum!

Hashing of Strings: Problems

HashString ("Tan Ah Teck")

$$= ("T" + "a" + "n" + " " + \\ "A" + "h" + " " + \\ "T" + "e" + "c" + "k") \% 11 \text{ // hash table size is 11}$$

$$= (84 + 97 + 110 + 32 + \\ 65 + 104 + 32 + \\ 84 + 101 + 99 + 107) \% 11$$

$$= 825 \% 11$$

$$= 0$$

Hashing of Strings: Problems

- Lee Chin Tan
- Chen Le Tian
- Chan Tin Lee

All 3 strings above have the **same hash value**! Why?

- **Problem:** The hash code produced does **not** depend on **positions** of characters! – Bad

Hashing Strings: Better Conversion

- Idea:

- Associate a **weight** to each position in string

- Common approach:

- Multiply each position by X^{position} , for a chosen X

- **Example:**

- Let's take $X = 17$
 - $\text{Hash}(\text{"abc"}) = 97 \cdot 17^2 + 98 \cdot 17^1 + 99 \cdot 17^0 = 29798$
 - Check whether "bac", "cba" etc gives different sum?

Hashing Strings: Better Conversion

- The idea can be implemented efficiently:
 - Using **Horner's Rule**

```
int HashString( string str )
{
    int sum = 0;

    for ( i = 0; i < str.size(); i++ )
        sum = (17*sum) + str[i];

    return sum % tableSizeM;
}
```

- In actual implementations, popular choice of X is 31 or 37

Hash Function: Summary

- First convert non-integer key into integer
 - Quality of conversion affects the hashing
- Perform hashing using the integer key
 - Take note of the range and characteristics of the input when designing hash function
 - Try to meet the qualities of a good hash function
- Modulo method is one of the most common choices for hash function

Collision Resolution

Probability of Collision (1/2)

- **von Mises Paradox (The Birthday Paradox)**: “How many people must be in a room before the probability that some **share a birthday**, ignoring the year and leap days, becomes at least 50 percent?”

$Q(n)$ = Probability of **unique** birthday for n people

$$= \frac{364}{365} \times \frac{363}{365} \times \frac{362}{365} \dots \frac{365 - n + 1}{365}$$

$P(n)$ = Probability of **collisions** (same birthday) for n people
 $= 1 - Q(n)$

$$P(\mathbf{23}) = \mathbf{0.507}$$

Probability of Collision (2/2)

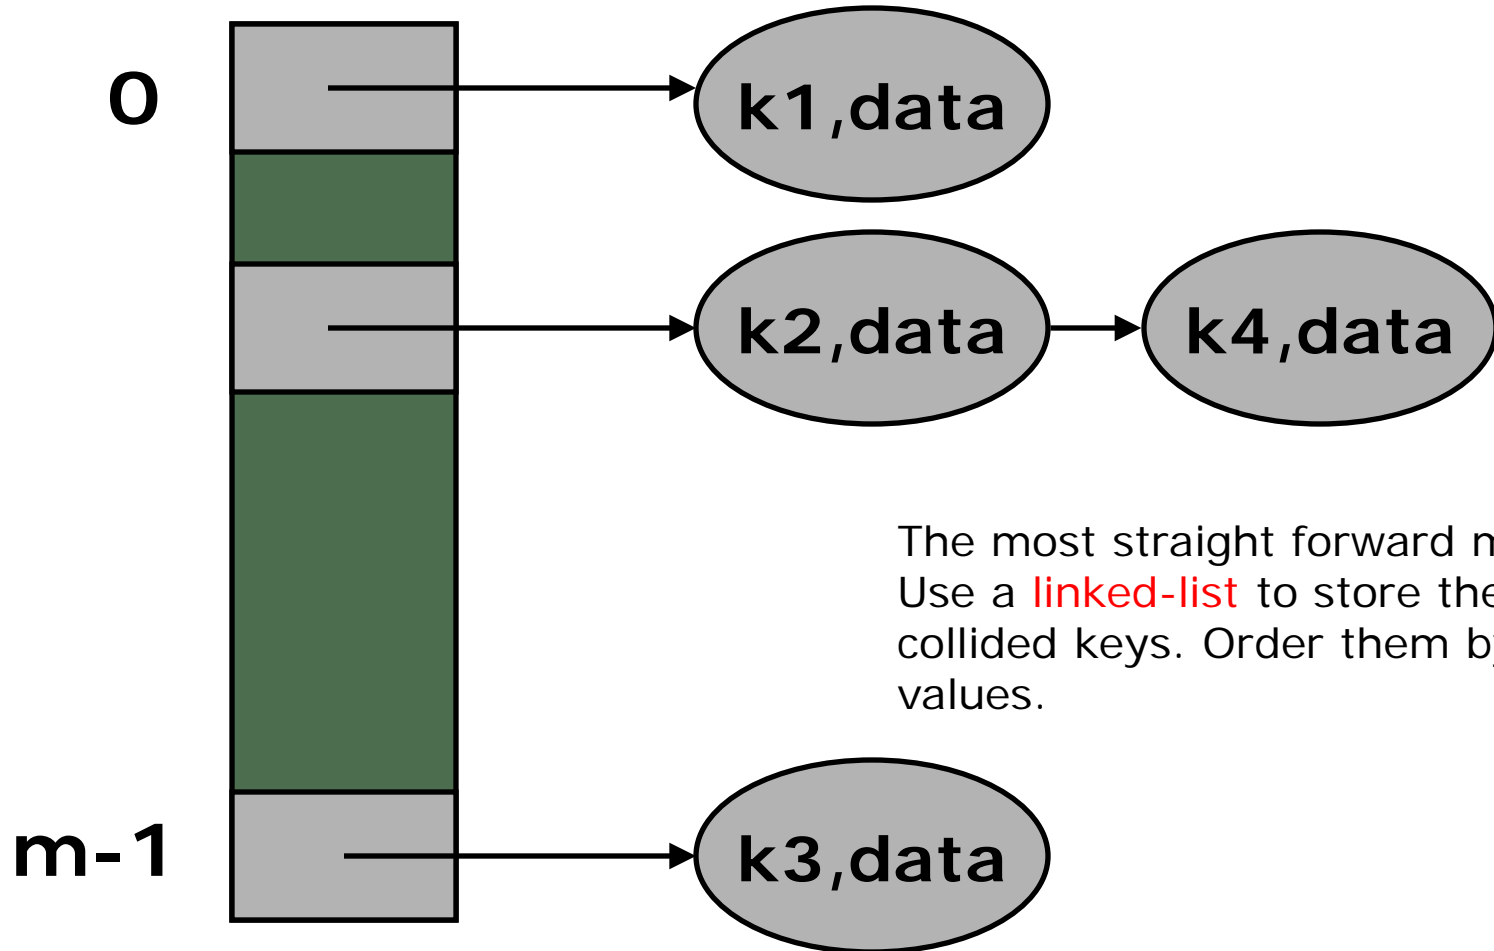
- This means that if there are **23** people in a room, the probability that some people share a birthday is **50.7%**.
- So, if we insert **23** keys in to a table with **365** slots, more than half of the time we get collisions. Such a result is counter-intuitive to many.
- So, collision is very likely!

How to resolve collisions?

Conflict resolution schemes commonly used:

- Separate Chaining
- Linear Probing
- Quadratic Probing
- Double Hashing

Separate chaining



The most straight forward method.
Use a **linked-list** to store the
collided keys. Order them by key
values.

Hash table (separate chaining)

insert (key, data)

insert data into the **list** $a[h(\text{key})]$

delete (key)

delete data from the **list** $a[h(\text{key})]$

find (key)

find key from the **list** $a[h(\text{key})]$

Load Factor

- **n**: number of keys in the hash table
- **m**: size of the hash table – number of slots
- Define the load factor α

$$\alpha = n/m$$

a measure of **how full** the hash table is.

(Note: α can be ≥ 1)

If table size is the number of linked lists, then α is the average length of the linked lists.

Average Running Time

- Find $O(1 + \alpha)$
 - Insert $O(1)$
 - Delete $O(1 + \alpha)$
-
- Note that α affects the performance of find and delete operations.
 - If α is bounded by some constant, then all three operations are $O(1)$.

Reconstructing hash table

- To keep α bounded, we may need to **reconstruct** the whole table when the load factor exceeds the bound.
- Whenever the load factor exceeds the bound, we need to **rehash** all keys into a **bigger** table (increase m to reduce α), say double the table size.

Open Addressing


- Separate chaining is a close addressing system as the address given to a key is fixed
- When the hash address given to a key is open (not fixed), the hashing is an open addressing system
- Open Addressing:
 - Hashed items are in a single array
 - Hash code gives the home address
 - Collision is resolved by checking multiple positions
 - Each check is called a probe into the table

Linear Probing

$$\text{hash}(k) = k \bmod 7$$

Here the table size $m=7$

Note: 7 is a prime number



In **linear probing**, when there is a **collision**, we scan forward for the the **next empty slot** (wrapping around when we reach the last slot)

Linear Probing: Insert 18

$$\text{hash}(k) = k \bmod 7$$

$$\begin{aligned}\text{hash}(18) \\ &= 18 \bmod 7 \\ &= 4\end{aligned}$$

0	
1	
2	
3	
4	18
5	
6	

Linear Probing: Insert 14

$$\text{hash}(k) = k \bmod 7$$

$$\begin{aligned}\text{hash}(14) \\ &= 14 \bmod 7 \\ &= 0\end{aligned}$$

0	14
1	
2	
3	
4	18
5	
6	

Linear Probing: Insert 21

$$\text{hash}(k) = k \bmod 7$$

$$\begin{aligned}\text{hash}(21) \\ &= 21 \bmod 7 \\ &= 0\end{aligned}$$

0	14
1	21
2	
3	
4	18
5	
6	



Collision occurs!
Look for **next empty slot**

Linear Probing: Insert 1

$$\text{hash}(k) = k \bmod 7$$

$$\begin{aligned}\text{hash}(1) \\ &= 1 \bmod 7 \\ &= 1\end{aligned}$$

0	14
1	21
2	1
3	
4	18
5	
6	

Collides with 21
(hash value 0). Look
for **next empty slot**

Linear Probing: Insert 35

$$\text{hash}(k) = k \bmod 7$$

$$\begin{aligned}\text{hash}(35) \\ &= 35 \bmod 7 \\ &= 0\end{aligned}$$

0	14
1	21
2	1
3	35
4	18
5	
6	

Collision, need to check **next 3 slots**

Linear Probing: Find 35

$$\text{hash}(k) = k \bmod 7$$

$$\text{hash}(35) = 0$$


0	14
1	21
2	1
3	35
4	18
5	
6	

Found 35, after 4 probes

Linear Probing: Find 8

$$\text{hash}(k) = k \bmod 7$$

$$\text{hash}(8) = 1$$

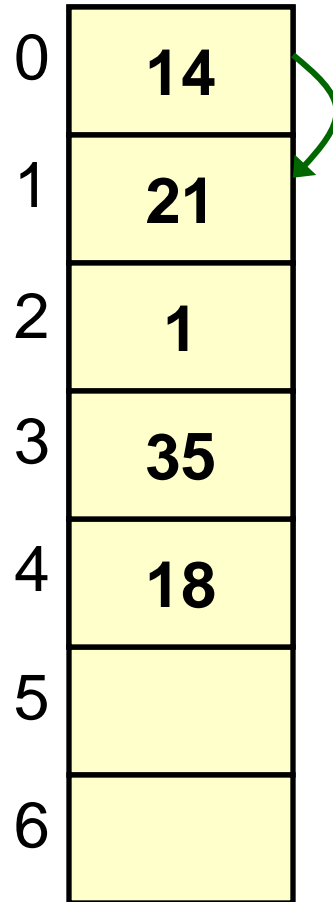


8 NOT found.
Need **5** probes!

Linear Probing: Delete 21

$\text{hash}(k) = k \bmod 7$

$\text{hash}(21) = 0$

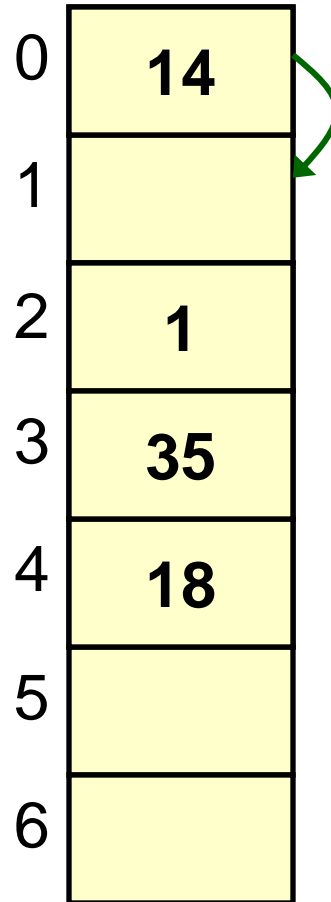


0	14
1	21
2	1
3	35
4	18
5	
6	

Linear Probing: Find 35

$\text{hash}(k) = k \bmod 7$

$\text{hash}(35) = 0$



0	14
1	
2	1
3	35
4	18
5	
6	

35 NOT found!
Incorrect!

We **cannot** simply **remove** a value, because it can affect **find()**!

How to delete?

- **Lazy** Deletion
- Use **three** different **states** of a slot
 - **Occupied**
 - **Deleted**
 - **Empty**
- When a value is removed from linear probed hash table, we just **mark** the status of the slot as “**deleted**”, Instead of emptying the slot
- Need to use a state array the same size as the hash table

Linear Probing: Delete 21

$$\text{hash}(k) = k \bmod 7$$

$$\text{hash}(21) = 0$$


0	14
1	21
2	1
3	35
4	18
5	
6	

Slot 1 is occupied but now **marked as deleted**.

Linear Probing: Find 35

$\text{hash}(k) = k \bmod 7$

$\text{hash}(35) = 0$



Found 35
Now we can find 35

Linear Probing: Insert 15

$$\text{hash}(k) = k \bmod 7$$

$$\text{hash}(15) = 1$$

Note: We **continue to search** for 15, and found that 15 is not in the hash table (total 5 probes).

0	14
1	15
2	1
3	35
4	18
5	
6	

15 is inserted into the slot 1 which was marked as deleted

We can insert a new value into a slot that has been marked as deleted.

Performance of Hash Table

Load factor	Number of Probes	
	Linear Probing	Chaining
0	1.00	1.00
0.25	1.17	1.13
0.5	1.50	1.25
0.75	2.50	1.38
0.83	3.38	1.43
0.9	5.50	1.45
0.95	10.50	1.48


For successful search, the number of probes is

Linear probing: $\frac{1}{2} (1 + 1/(1 - \alpha))$

Separate chaining: $1 + \alpha/2$

Problem of Linear Probing

- A cluster is a collection of consecutive occupied slots
- A cluster that covers the home address of a key is called a primary cluster of the key
- Linear probing can create large primary clusters that will increase the running time of find/insert/delete operations



Problem of Linear Probing

The probe sequence of linear probing is:

$\text{hash}(\text{key})$	// first probe, the home
$(\text{hash}(\text{key}) + 1) \% m$	// second probe
$(\text{hash}(\text{key}) + 2) \% m$	// third probe
$(\text{hash}(\text{key}) + 3) \% m$	// fourth probe

- If there is an empty slot, we are sure to find it.
- When an empty slot is found, conflict is resolved, but the primary cluster of the key is expanded as a result
- The size of the resulting primary cluster may be very big due to the annexation of the neighbouring cluster

Modified Linear Probing

Q: How to modify linear probing to **avoid primary clustering**?

We can modify the **probe sequence** as follows:

$$\begin{aligned} & \text{hash(key)} \\ & (\text{hash(key)} + \mathbf{1} * \mathbf{d}) \% m \\ & (\text{hash(key)} + \mathbf{2} * \mathbf{d}) \% m \\ & (\text{hash(key)} + \mathbf{3} * \mathbf{d}) \% m \\ & \vdots \end{aligned}$$

where **d** is some constant integer >1 and is co-prime to **m**.

Note: Since **d** and **m** are co-primes, the probe sequence **covers all** the slots in the hash table.

Quadratic Probing

To escape from the primary cluster quickly,
use quadratic probing to look for an empty
slot.

The probe sequence is

hash(key)	
(hash(key) + 1) % m	jump 1²
(hash(key) + 4) % m	jump 2²
(hash(key) + 9) % m	jump 3²
:	

Distance from
previous probe

+1

+3

+5

+7

...

**+2j-1 for jth probe
from home**

Quadratic Probing: Insert 3, 18

hash(k) = k mod 7


hash(3) = 3

hash(18) = 4

0	
1	
2	
3	3
4	18
5	
6	

Quadratic Probing: Insert 38

$$\text{hash}(k) = k \bmod 7$$



$$\text{hash}(38) = 3$$

Collision

Can quadratic probing always find a free slot?

Insert 12 into the previous example, followed by 10.

what happens?

Theorem

If $\alpha < 0.5$, and m is prime, then we can always find an empty slot.

(m is the table size and α is the load factor)

When quadratic probing is used, in the worst case, 50% of the hash table is wasted.

Secondary Clustering

- In quadratic probing, clusters are formed along the path of probing, instead of around the home location
- These clusters are called **secondary clusters**
- Secondary clusters are formed as a result of using the same pattern in probing by all keys

Double Hashing

- To resolve the secondary clustering problem, we have to break the probing pattern of quadratic hashing
- We may use another hash function hash_2 to generate different probe sequences for different keys

$\text{hash}(\text{key})$

$(\text{hash}(\text{key}) + 1 * \text{hash}_2(\text{key})) \% m$

$(\text{hash}(\text{key}) + 2 * \text{hash}_2(\text{key})) \% m$


$(\text{hash}(\text{key}) + 3 * \text{hash}_2(\text{key})) \% m$

:

hash_2 is called the **secondary hash function**, the no of slots to jump each time a collision occurs

After Inserting 14 and 18, Insert 21

$\text{hash}(k) = k \bmod 7$
 $\text{hash}_2(k) = k \bmod 5$




$\text{hash}(21)$
 $= 21 \bmod 7$
 $= 0$

$\text{hash}_2(21)$
 $= 21 \bmod 5$
 $= 1$

Double Hashing: Insert 4

$\text{hash}(k) = k \bmod 7$
 $\text{hash}_2(k) = k \bmod 5$




If we insert 4, the
probe sequence is
4 (home), 8, 12, ...

$\text{hash}(4) = 4$
 $\text{hash}_2(4) = 4$

Double Hashing: Insert 35

$\text{hash}(k) = k \bmod 7$
 $\text{hash}_2(k) = k \bmod 5$

$\text{hash}(35) = 0$
 $\text{hash}_2(35) = 0$



But if we insert 35,
the probe sequence
is **0, 0, 0, ...**

What is wrong?
Since $\text{hash}_2(35) = \mathbf{0}$.
Not acceptable!

Warning

- **Secondary hash function must **not** evaluate to **0** !**
- To solve this problem, simply change $\text{hash}_2(\text{key})$ in the above example to:

$$\text{hash}_2(\text{key}) = 5 - (\text{key} \% 5)$$

Note: If $\text{hash}_2(k) = 1$, then it is the same as linear probing.

If $\text{hash}_2(k) = d$, where d is a constant and $d > 1$, then it is the same as modified linear probing.

Good Collision Resolution Method

- Small cluster size
- Always find an empty slot if it exists
- Give different probe sequences when 2 keys collide (i.e. no secondary clustering)
- Fast

Rehash

- Time to rehash:
 - When table is getting full, the operations are getting slow
 - For quadratic probing, inserts might fail when the table is more than half full
- Rehash operation:
 - Build another table about twice as big with a new hash function
 - Scan the original table, for each key, compute the new hash value and insert the data into the new table
 - Delete the original table
- The load factor used to decide the time to rehash:
 - For open addressing: 0.5
 - For closed addressing: 1

STL unordered_map

- STL unordered_map implements a Hash Table with separate chaining:
- Associate containers that store elements formed by the combination of a key value and a mapped value, and which allows for fast retrieval of individual elements based on their keys
- https://en.cppreference.com/w/cpp/container/unordered_map

unordered_map example

```
// std::unordered_map
#include <bits/stdc++.h>

int main()
{
    // Unordered map
    std::unordered_map<int, int> order;

    // Mapping values to keys
    order[5] = 10;
    order[3] = 5;
    order[20] = 100;
    order[1] = 1;

    // Iterating the map and printing unordered values
    for (auto i = order.begin(); i != order.end(); i++) {
        std::cout << i->first << " : " << i->second << "\n";
    }
}
```

Output:

```
1 : 1
3 : 5
20 : 100
5 : 10
```


STL unordered_set

- There is also an unordered_set if key->value pairs are not required
- Implemented using hash table where keys are stored in any order
- https://en.cppreference.com/w/cpp/container/unordered_set

unordered_set example

```
// C++ program to demonstrate various function of unordered_set
#include <bits/stdc++.h>
using namespace std;

int main()
{
    // declaring set for storing string data-type
    unordered_set<string> stringSet;

    // inserting various string, same string will be stored
    // once in set
    stringSet.insert("code");
    stringSet.insert("in");
    stringSet.insert("c++");
    stringSet.insert("is");
    stringSet.insert("fast");

    string key = "slow";
```

Example (cont'd)

```
// find returns end iterator if key is not found,
// else it returns iterator to that key
if (stringSet.find(key) == stringSet.end())
    cout << key << " not found\n\n";
else
    cout << "Found " << key << endl << endl;

key = "c++";
if (stringSet.find(key) == stringSet.end())
    cout << key << " not found\n";
else
    cout << "Found " << key << endl;

// now iterating over whole set and printing its
// content
cout << "\nAll elements : ";
unordered_set<string> :: iterator itr;
for (itr = stringSet.begin(); itr != stringSet.end(); itr++)
    cout << (*itr) << endl;
}
```

Output:
slow not found

Found c++

All elements :
is
fast
c++
in
code

Summary

- How to hash? Criteria for good hash functions?

- How to **resolve collision**?

Collision resolution techniques:

- separate chaining
- linear probing
- quadratic probing
- double hashing

- Problem on deletions

- **Primary** clustering and **secondary** clustering

- STL unordered_map, unordered_set