
CS2040C Data Structures and Algorithms

Single-source Shortest Path


(Bellman-Ford's Algorithm)

Outline

- Definitions
- Unweighted Shortest Paths
- SSSP for positive weighted graphs
- Relax(u, v)
- Bellman-Ford's Algorithm
- Running Time of Bellman-Ford
- Special Case/s
 - BFS for SSSP

Definitions

- A **path** on a graph G is a sequence of vertices $v_0, v_1, v_2, \dots, v_n$ where $(v_i, v_{i+1}) \in E$
- The **cost** of a path is the **sum** of the cost of all edges in the path.



In the **single-source shortest path (SSSP)** problem, we are given a vertex s , and we want to find the path with minimum cost (weight) to **every other vertex**.

Definitions

distance(v): shortest distance **so far** from **s** to v

parent(v): previous node on the shortest path so far from s to v

weight(u, v): the weight (**cost**) of edge from u to v


$\delta(u, v)$: actual shortest path from u to v

Note: $\text{weight}(s, s) = 0$

$\text{weight}(s, u)$ where u is unreachable $= +\infty$

Unweighted shortest path


If a graph is **unweighted**, we can treat the **cost** of each edge as **1**.



ShortestPath(s)

- The shortest path for an *unweighted* graph can be found using BFS.
- Run BFS(**s**) where **s** is the chosen **source** node
- Trace back the parent pointer from v to s to get the shortest path
- no. of edges in the path is given by the level of a vertex in the BFS tree (or level – 1, if level of root is 1)


ShortestPath(s)




Question: Why does BFS guarantee shortest paths?

Positive weighted shortest path

Will BFS work?



Weighted graphs




e.g.: weight = distance

Shortest paths

Questions:


- How far is it from S to D?
- What is the shortest path from S to D?
- Find the shortest path from S to every node
- Find the shortest path between every pair of nodes




Will BFS work?

Cannot use **BFS**


BFS finds minimum number of **HOPS** not minimum **DISTANCE**




Distance from the source?




Distance from the source?




An example: BFS




An example: BFS



An example: BFS




An example: BFS




An example: BFS

BFS finds minimum number of **HOPS** not minimum **DISTANCE**




BFS(s) does **not** work



- Must keep track of **shortest** distance from the source node **so far** for each node
- **Observation 1:** If we found a new shorter path, update the distance.


Observation 1

- In the following figures, we label a node with the shortest distance discovered so far from the source.
- Here is the basic idea that will help us solve our shortest path problem.
- If the current shortest distance from s to v is 13, to u is 10, and the cost of edge (u,v) is 1, then we have discovered a shorter path from s to v (through u).




Observation 2 (for positive costs only)

- The **second idea** is that **if** we know the shortest distance so far from s to v is 6, and the shortest distance so far from s to other nodes (in white) is **bigger or equal** to 6, **then** there **cannot** be a shorter path to v through these other nodes.
- This is **true only if** the costs are **positive**!



Example



$\text{distance}(v) = 8$

$\text{weight}(u, v) = 2$

$\text{parent}(v) = u$


Relax(u, v) - based on observation 1

$d = \text{distance}(u) + \text{weight}(u, v)$

if $\text{distance}(v) > d$ **then** // found a **new shorter distance**

$\text{distance}(v) = d$ // update the **distance** and **parent**

$\text{parent}(v) = u$



Data structures needed

- Array/Vector **dist** of size **V** (dist: distance)
Initially **dist[u]** = 0 if $u = s$; else **D[u]** = $+\infty$ (10^9)
 - **dist[u]** decreases as we find better (shorter) paths
 - **dist[u]** $\geq \delta(s, u)$ throughout the execution of SSSP algorithm
 - **dist[u]** = $\delta(s, u)$ at the end of SSSP algorithm
- Array/Vector **p** of size **V** (p: parent/predecessor)
 - **p[u]** = the predecessor on best path from source **s** to **u**
 - **p[u]** = NULL (initially not defined, we can use -1)

This array/Vector **p** describes the resulting SSSP spanning tree

shortest paths

Maintain estimate for each distance:


- ❑ **Reduce** estimate
- ❑ **Invariant:** estimate \geq shortest distance

The idea:

$\text{relax}(u, v)$:

Test if the best way to get from $s \rightarrow v$ is to go from $s \rightarrow u$, then $u \rightarrow v$.

Update dist




shortest paths

Maintain estimate for each distance:

$\text{relax}(S, A)$

```
relax(int u, int v){  
    if (dist[v] > dist[u] +  
        weight(u,v))  
        dist[v] = dist[u] +  
            weight(u,v);  
    p[v] = u;  
}
```




shortest paths

Maintain estimate for each distance:

$\text{relax}(S, A)$

```
relax(int u, int v){  
    if (dist[v] > dist[u] +  
        weight(u,v))  
        dist[v] = dist[u] +  
            weight(u,v);  
    p[v] = u;  
}
```




shortest paths

Maintain estimate for each distance:

$\text{relax}(A, C)$

```
relax(int u, int v){  
    if (dist[v] > dist[u] +  
        weight(u,v))  
        dist[v] = dist[u] +  
            weight(u,v);  
    p[v] = u;  
}
```




shortest paths

Maintain estimate for each distance:

`relax(A, C)`

```
relax(int u, int v){  
    if (dist[v] > dist[u] +  
        weight(u,v))  
        dist[v] = dist[u] +  
            weight(u,v);  
    p[v] = u;  
}
```




shortest paths

Maintain estimate for each distance:

`relax(A, B)`

```
relax(int u, int v){  
    if (dist[v] > dist[u] +  
        weight(u,v))  
        dist[v] = dist[u] +  
            weight(u,v);  
    p[v] = u;  
}
```




shortest paths

Maintain estimate for each distance:

$\text{relax}(S, B)$

```
relax(int u, int v){  
    if (dist[v] > dist[u] +  
        weight(u,v))  
        dist[v] = dist[u] +  
            weight(u,v);  
    p[v] = u;  
}
```




shortest paths

Maintain estimate for each distance:

`relax(B, C)`


```
relax(int u, int v){  
    if (dist[v] > dist[u] +  
        weight(u,v))  
        dist[v] = dist[u] +  
            weight(u,v);  
    p[v] = u;  
}
```



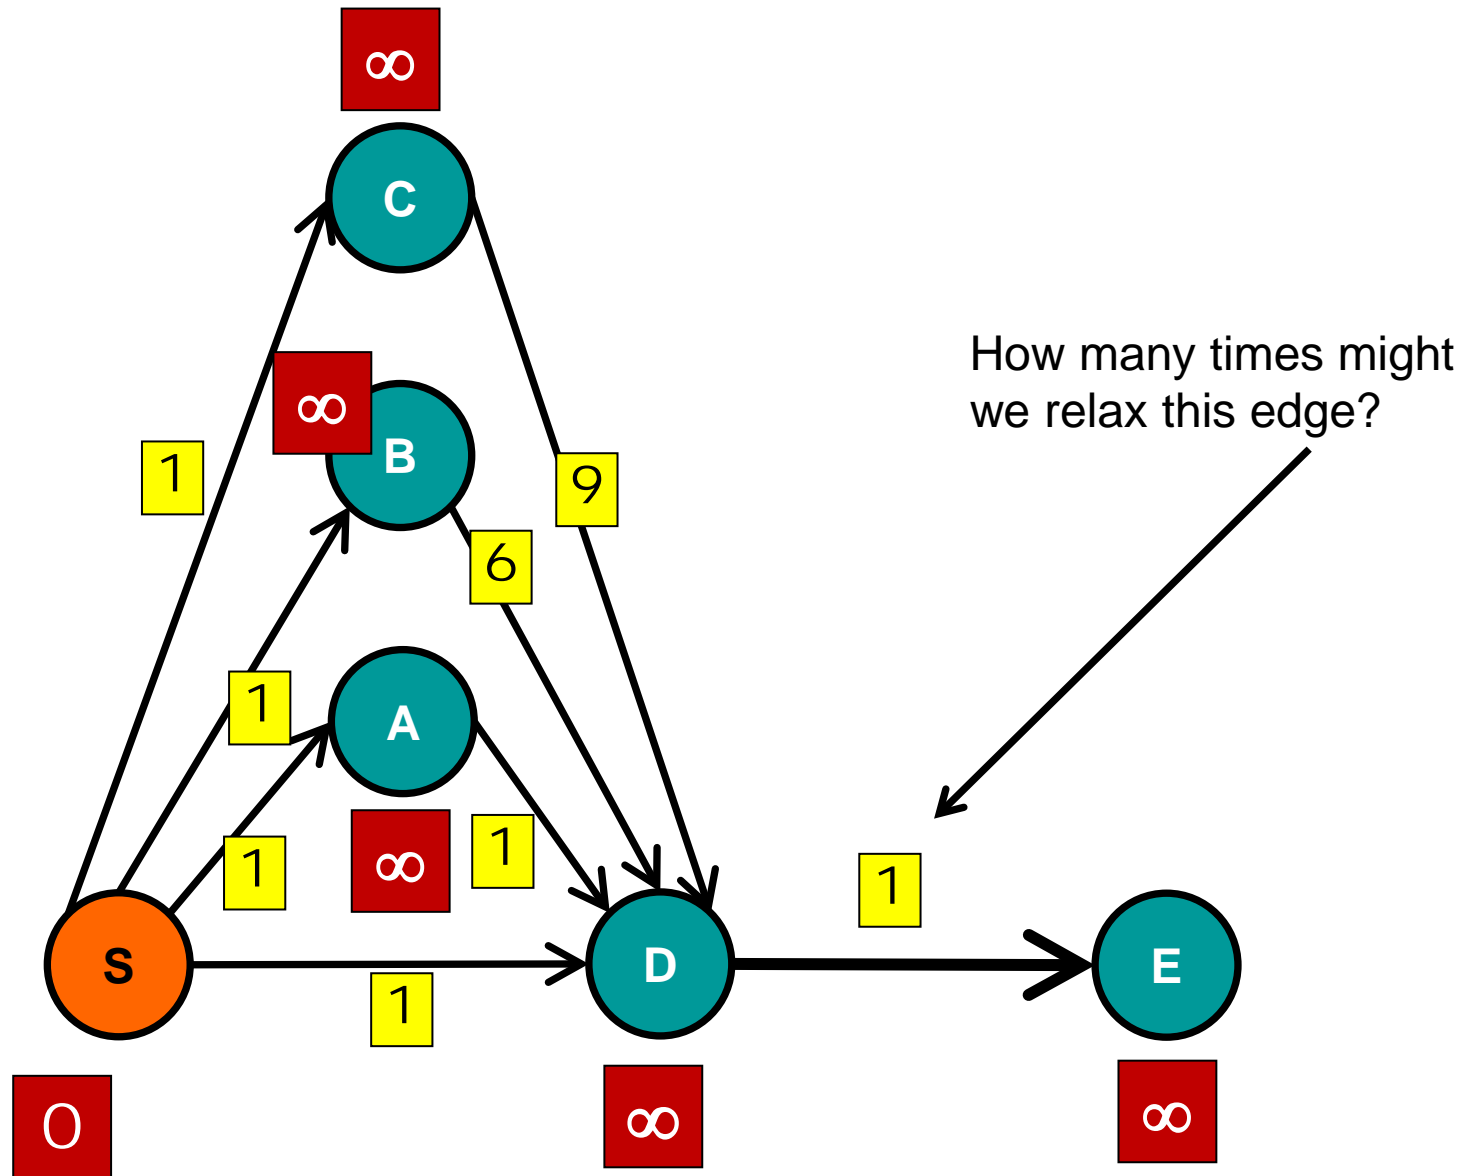
shortest paths

Maintain estimate for each distance:

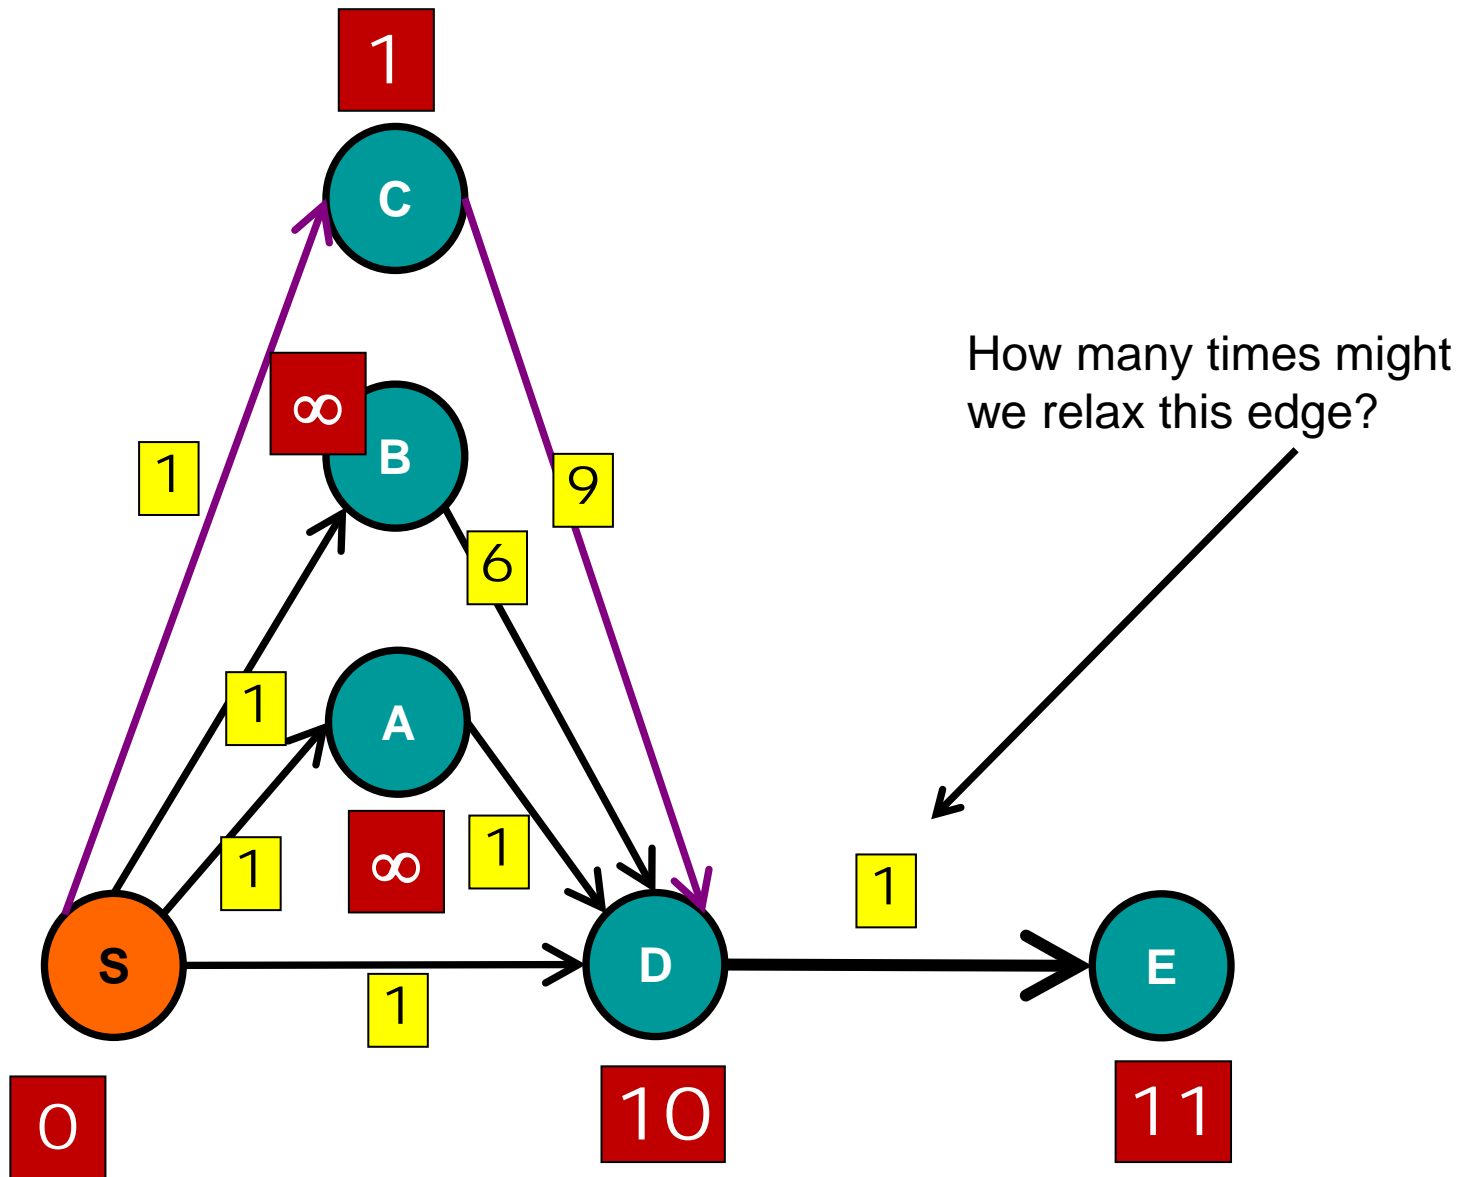
```
for Edge e in  
graph  
    relax(e)
```



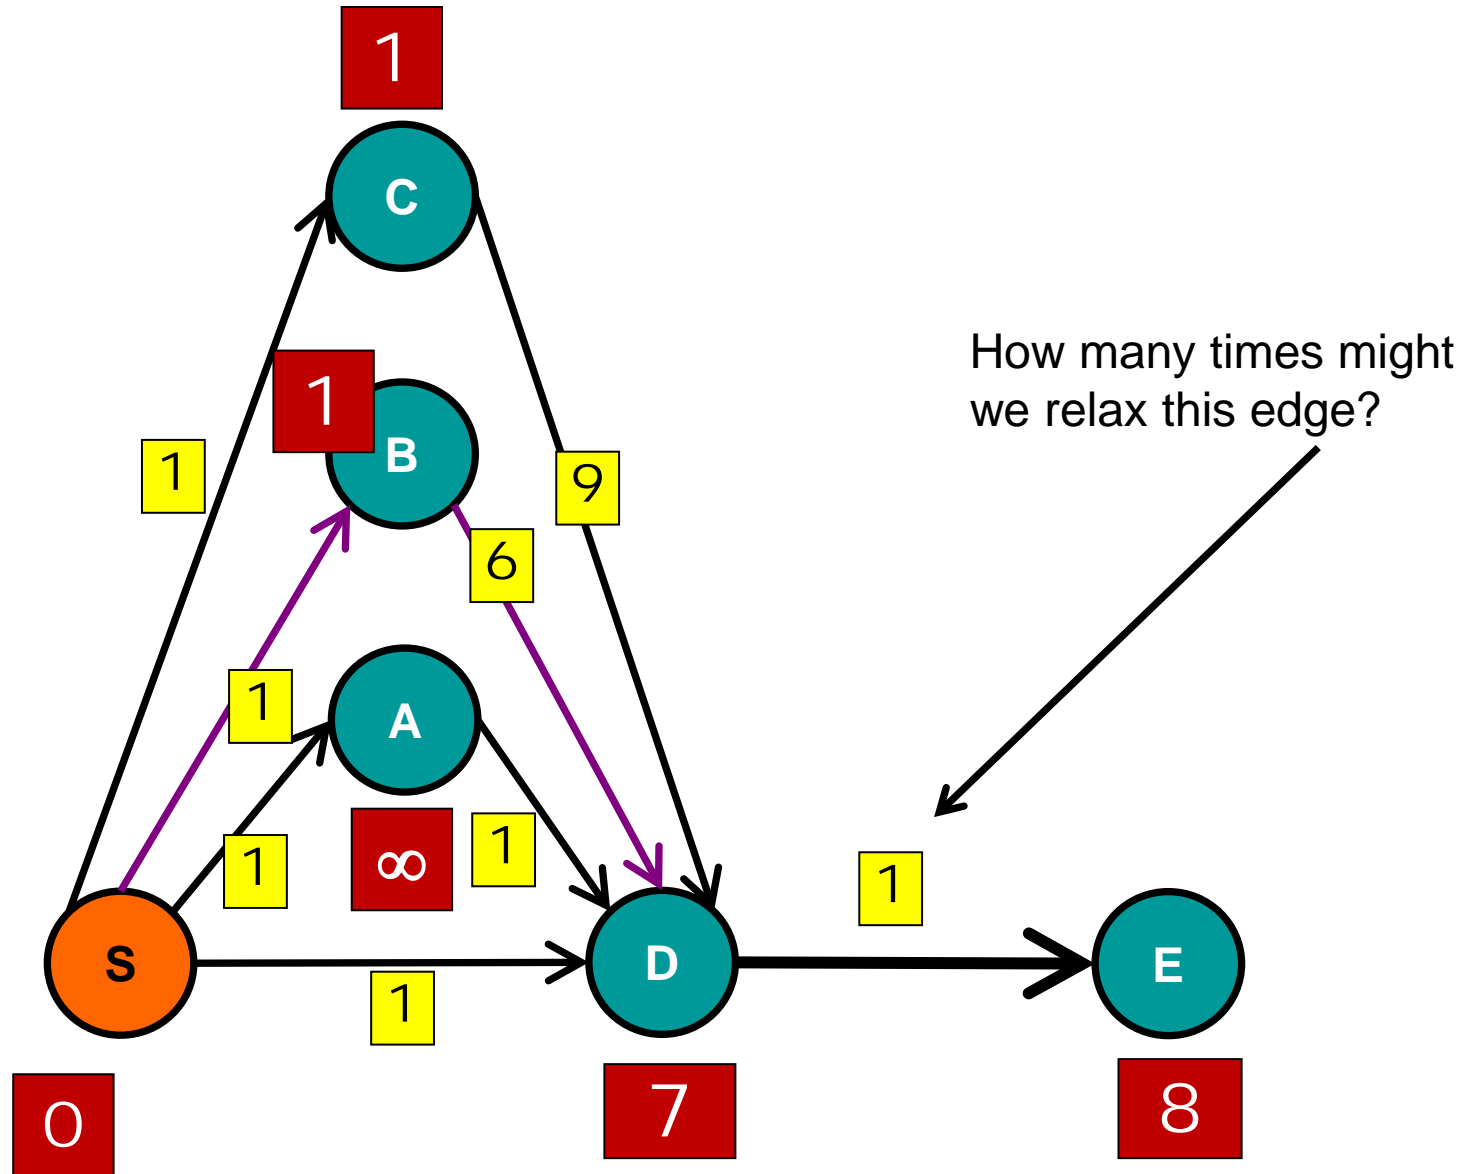
relaxation



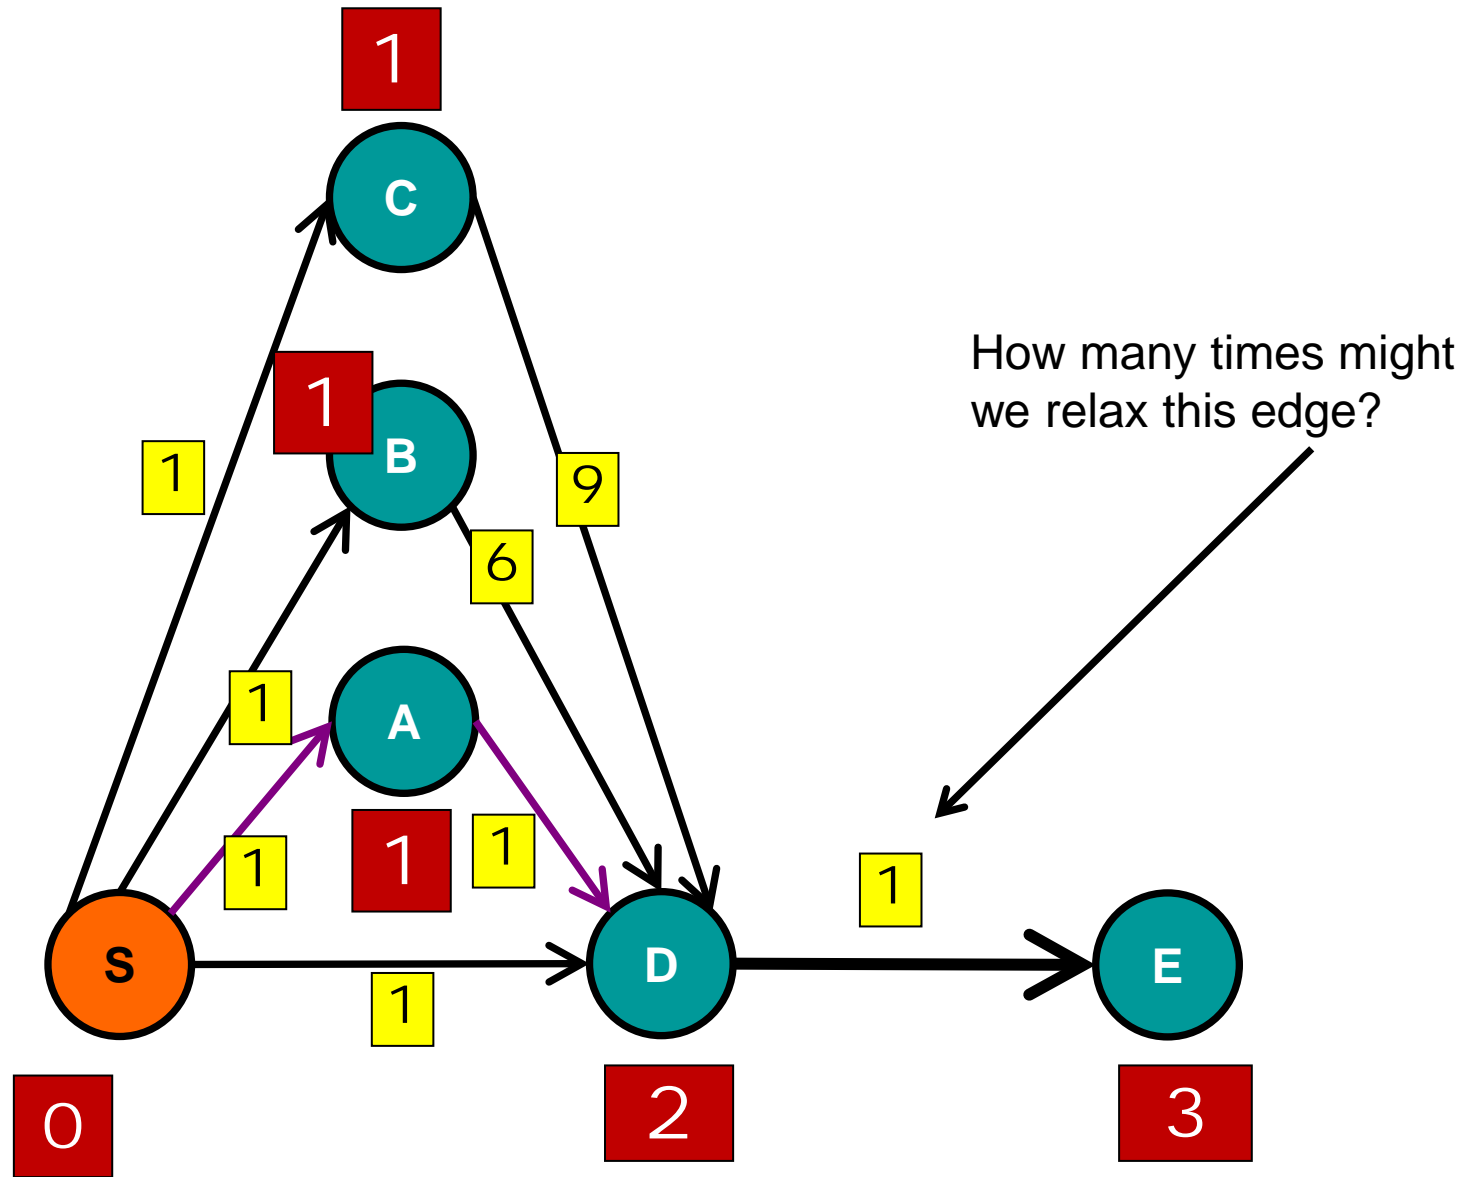
relaxation



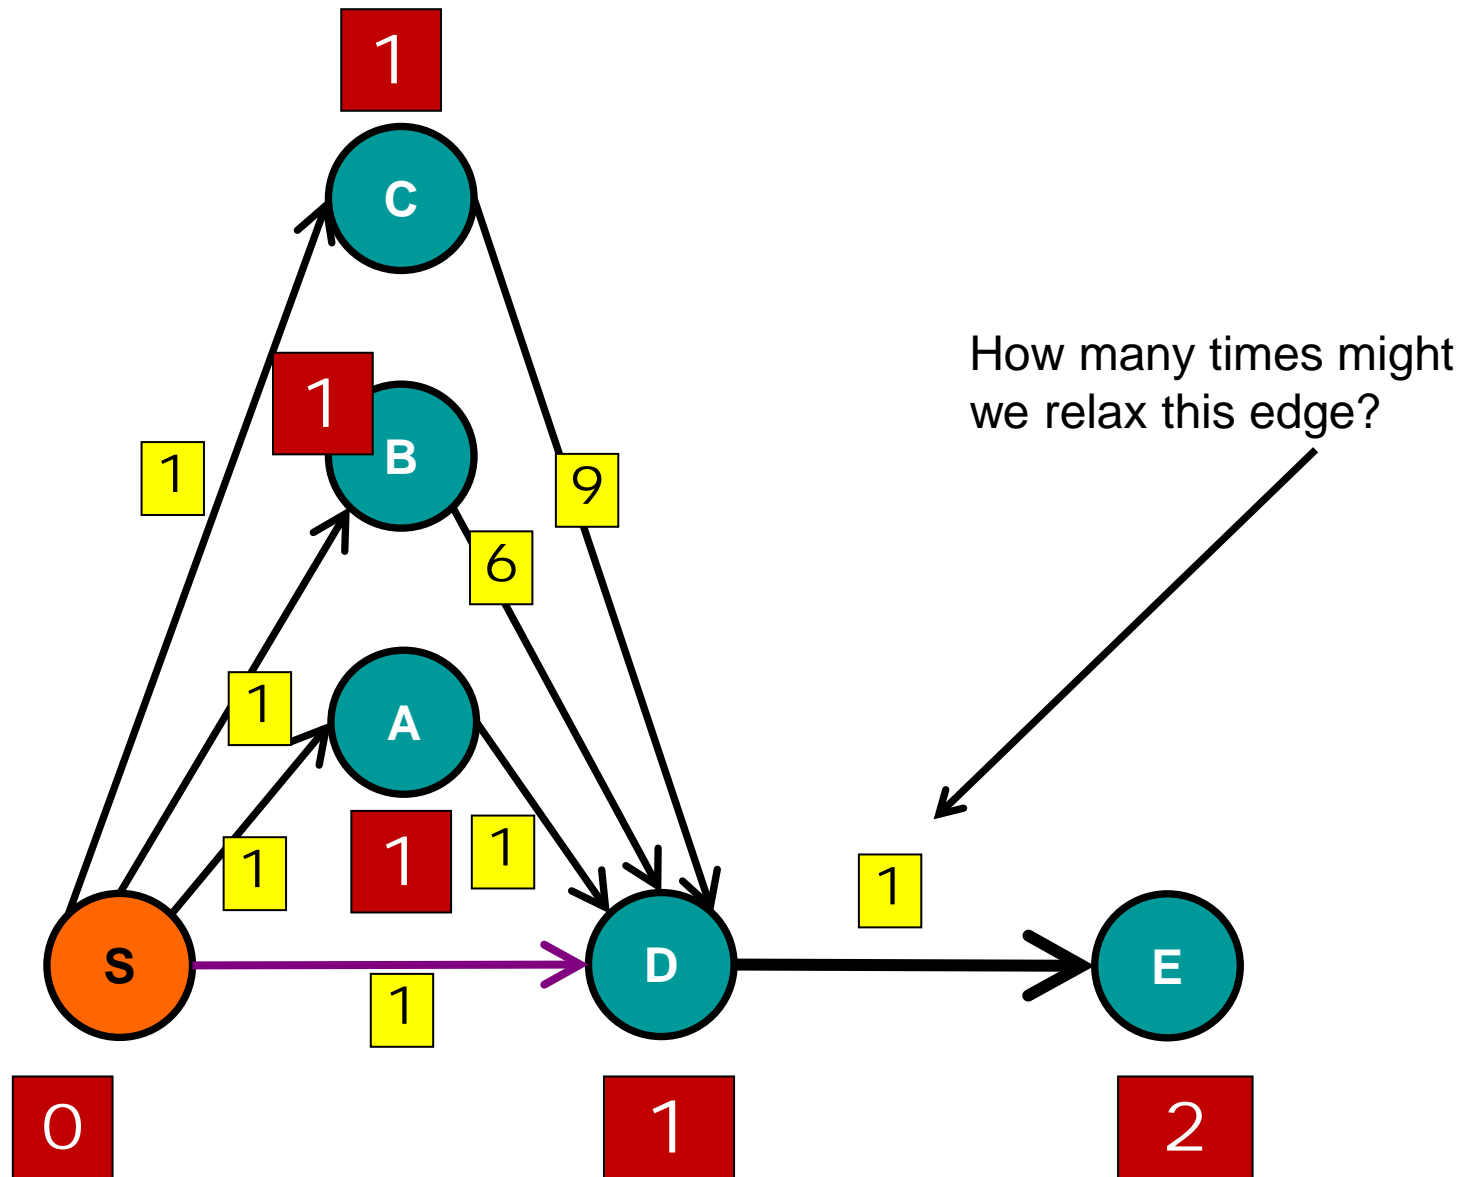
relaxation



relaxation




relaxation



Bellman-Ford algorithm

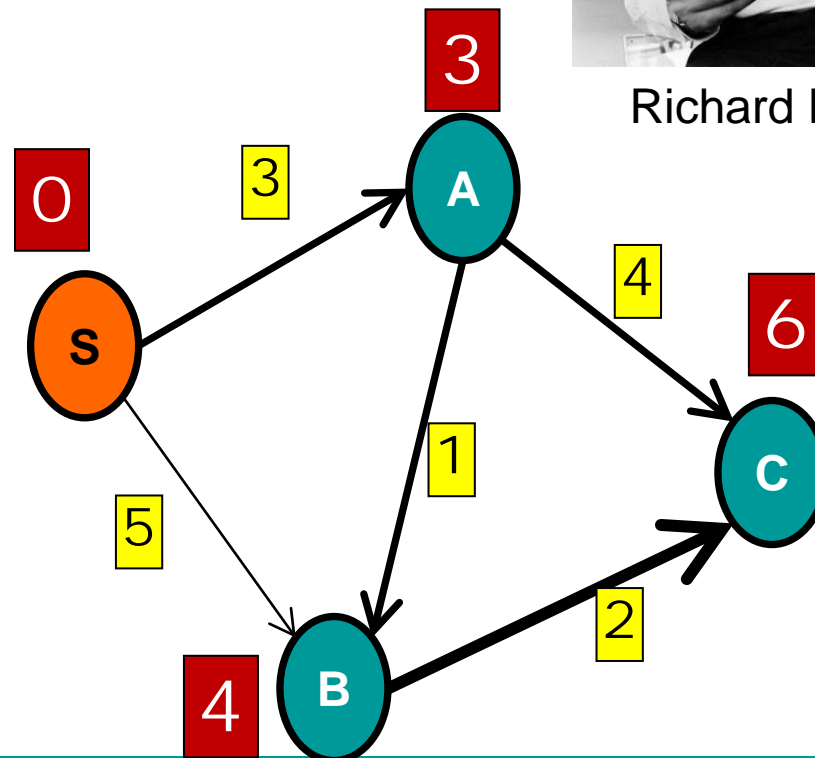
```
n = V.length
for i = 1 to n-1
  for Edge e in Graph
    relax(e)
```




Richard Bellman

**Does Bellman-Ford
always work?**


**Yes! Proof by
Induction (in
Visualgo)**



why does Bellman-Ford work?




why does Bellman-Ford work?



Look at minimum weight path from S to D.
(Path is simple: no loops)


why does Bellman-Ford work?



After 1 iteration, 1 hop estimate is correct.


meaning: All shortest paths that are 1 hop long are now correct

why does Bellman-Ford work?




After 1 iteration, 1 hop estimate is correct.

why does Bellman-Ford work?




After 1 iteration, 1 hop estimate is correct.

why does Bellman-Ford work?




After 2 iterations, 2 hop estimate is correct.

why does Bellman-Ford work?




After 3 iterations, 3 hop estimate is correct.

why does Bellman-Ford work?



After 4 iterations, D estimate is correct.

Why does Bellman-Ford work?




Keep running till $V-1$ and Bellman-Ford finds shortest paths from s to all other nodes!

What is the running time of Bellman-Ford?

```
n = V.length
for i = 1 to n-1
    for Edge e in Graph
        relax(e)
```


What is the running time of Bellman-Ford?

- A. $O(V)$
- B. $O(E)$
- C. $O(V + E)$
-  D. $O(VE)$
- E. $O(E \log V)$
- F. I have no idea.

Early termination?

```
n = V.length
for i = 1 to n-1
    for Edge e in Graph
        relax(e)
```

When can we terminate early?


- A. When a relax operation has no effect.
- B. When two consecutive relax operations have no effect.
-  C. When an entire sequence of $|E|$ relax operations have no effect.
- D. Never. Only after $|V|$ complete iterations.

Shortest paths

Maintain estimate for each distance:

```
for Edge e in  
graph  
    relax(e)
```

If we relax all the edges and there is no faster way to get to any node, we have the shortest paths!




Negative edge weights?


Will Bellman-Ford algorithm still work?

**Bellman-Ford
has no
problems with
negative edge
weights!**


....almost....



What if the graph looks like this?



negative weight **cycle**




$d(S,C)$ is infinitely negative!

negative weight **cycle**


Run Bellman-Ford for $|V|$ iterations.

If an **estimate** changes in the **last iteration** then negative weight cycle.



How to **detect** negative weight cycles?

Special case:



all edges have the **same weight**: What can we use?

BFS!!!

BFS for SSSP

We need to perform some simple modifications to BFS for it to be able to solve the unweighted version (or equal weights) of the SSSP problem:

1. we change the Boolean array **visited** into an Integer array **D**.
2. At the start of BFS, instead of setting `visited[u] = false`, we set `D[u] = 1e9` (a large number to symbolise $+\infty$ or even -1 to symbolise 'unvisited' state, but we cannot use 0 as `D[0] = 0`) $\forall u \in V \setminus \{s\}$; Then we set `D[s] = 0`
3. We change the BFS main loop from
if (`visited[v] == 0`) { `visited[v] = 1` ... } // v is unvisited
to
if (`D[v] == 1e9`) { `D[v] = D[u] + 1` ... } // v is 1 step (or whatever
// weight) away from u

Other special cases

Condition	Algorithm	Time Complexity
No Negative Weight Cycles	Bellman-Ford Algorithm	$O(VE)$
On Unweighted Graph (or equal weights)	BFS	$O(V + E)$
No Negative Weights	Dijkstra's Algorithm	$O((V + E) \log V)$
On Tree	BFS / DFS	$O(V)$
On DAG	Dynamic Programming	$O(V + E)$

Summary

- Single source shortest path (very common Computer Science problem)
- SSSP: Given source vertex s , want to find shortest path weights path to every other vertex v
- Bellman-Ford can be used in graphs that contain negative weights, but not negative weight cycles
- BFS can be used for unweighted or paths with same weights

*Acknowledgement: some slides courtesy of Dr Harold Soh

Programming Exam

- Date: 3 Nov 2018 (week 11)
- Time: 1pm to 3pm
- Venues: PL1 and PL2
- Lab allocation has been uploaded in IVLE
- Format: 2 problems each with subtasks
- Total 12%
- Open Book (allowed to bring hard copy material)
- No internet access available during PE (i.e. cannot use online IDE or sunfire)