
CS2040C: Data Structures and Algorithms

Single Source Shortest Paths
(more special cases)

Outline

SSSP for special cases and the algorithms that are applicable and can run faster for these cases


- Using BFS/DFS on Trees
- Dijkstra's algorithm for graphs with no negative weights
- Modified Dijkstra's algorithm for graphs with negative weights
- Dynamic programming for DAGs

Special Cases

We have already covered the first two cases in the previous lecture

Condition	Algorithm	Time Complexity
No Negative Weight Cycles	Bellman-Ford Algorithm	$O(VE)$
On Unweighted Graph (or equal weights)	BFS	$O(V + E)$
No Negative Weights	Dijkstra's Algorithm	
Negative weights	Modified Dijkstra's Algorithm	
On Tree	BFS / DFS	
On DAG	Dynamic Programming (one-pass Bellman-Ford)	

Special Case: Undirected, Weighted Tree



Trees (redefined)

What is an (undirected) tree?

- ❑ A graph with no cycles is an (undirected) tree

What is a *rooted* tree?


- ❑ A tree with a special designated root node

Our previous (recursive) definition of a *tree*:

- ❑ A node with zero, one, or more sub-trees
- ❑ a *rooted* tree


Undirected Weighted Tree

Assume you can only cross an edge once on your path.




Undirected Weighted Tree

**how many ways to
get from S to D?
(assume no
backpedalling)**



Undirected Weighted Tree


Just 1 way! It's a tree!



Tree: source-to-all


**In what order
should we relax the
nodes?**

Use DFS or BFS




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order




Tree: source-to-all

Relax in DFS Order



Why does this work?

Tree: source-to-all




Once you update a node, you never have to update it again!

Undirected Weighted Tree

Time Complexity?


every node only has one parent (except the root).
 $O(V) = O(E)$ edges.



Special Cases


Condition	Algorithm	Time Complexity
No Negative Weight Cycles	Bellman-Ford Algorithm	$O(VE)$
On Unweighted Graph (or equal weights)	BFS	$O(V + E)$
→ No Negative Weights	Dijkstra's Algorithm	
Negative Weights	Modified Dijkstra's Algorithm	
On Tree	BFS / DFS	$O(V)$
On DAG	Dynamic Programming	

General graph: non-negative edges



General graph: non-negative edges


Shortest paths form a **tree**



General graph: non-negative edges

Key property:

If p is the shortest path from S to D ,
and if p goes through X ,
then p is also the shortest path from S to X (and from X to D).



Dijkstra's Algorithm

Key idea:


Relax the edges in the “right” order.

Only relax each edge **once**:

- $O(E)$ cost (for relaxation step)

Edsger W. Dijkstra


- *“Computer science is no more about computers than astronomy is about telescopes.”*
- *“The question of whether a computer can think is no more interesting than the question of whether a submarine can swim.”*
- *“There should be no such thing as boring mathematics.”*
- *“Elegance is not a dispensable luxury but a factor that decides between success and failure.”*
- *“Simplicity is prerequisite for reliability.”*



1930-2002

Edsger W. Dijkstra

- *“It is practically impossible to teach good programming to students that have had a prior exposure to BASIC: as potential programmers they are mentally mutilated beyond hope of regeneration.”*
- *“The use of COBOL cripples the mind; its teaching should, therefore, be regarded as a criminal offense.”*
- *“Object-oriented programming is an exceptionally bad idea which could only have originated in California.”*




1930-2002

Edsger W. Dijkstra

From Wikipedia:

- *His approach to teaching was unconventional ...*
- *He invited the students to suggest ideas, which he then explored, or refused to explore because they violated some of his tenets.*
- *He conducted his final examinations orally, over a whole week.*
- *Each student was examined in Dijkstra's office or home, and an exam lasted several hours.*




1930-2002

Dijkstra's algorithm

Basic idea:

- ❑ Maintain distance estimate for every node
- ❑ Begin with empty shortest-path-tree
- ❑ Repeat:
 - Consider vertex with minimum **estimate**
 - Add vertex to **shortest-path-tree**
 - Relax all outgoing edges

Dijkstra's algorithm




Dijkstra's algorithm


Use node **A** as the **source**
i.e. find shortest distances
of all nodes from source
node A

The shortest
distance of A from A
is **0**


Initialize the
distances of others
nodes as **infinite**




Dijkstra's algorithm



Dijkstra's algorithm




Dijkstra's algorithm




(1) Update
distance for 3
nodes

Dijkstra's algorithm




(2) Select the next node with smallest distance value


Dijkstra's algorithm



Dijkstra's algorithm




Dijkstra's algorithm




Dijkstra's algorithm

The distances of all nodes now are the **shortest distances** from the source node **A**



Dijkstra's algorithm

The **shortest paths** from
the source node **A**



Dijkstra's algorithm

color all vertices **yellow**

// yellow nodes are those not yet processed

foreach vertex w

distance(w) = INFINITY

distance(s) = 0 //source node distance is 0

Dijkstra's algorithm


while there are **yellow** vertices //unprocessed nodes are yellow

v = **yellow** vertex with **min distance(v)**

color **v** **red** // red vertices are vertices with shortest distances from s found

foreach **yellow neighbour** **w** of **v**

relax(v,w)



Time Complexity

color all vertices yellow

foreach vertex w

$\text{distance}(w) = \text{INFINITY}$

$\text{distance}(s) = 0$

while there are yellow vertices

$v = \text{yellow vertex with min distance}(v)$

 color v red

foreach yellow neighbour w of v

$\text{relax}(v, w)$

Time Complexity

- Initialization takes $O(V)$ time. // V = no of nodes
- Picking the vertex with **minimum** distance(v) can take **$O(V)$** time, and relaxing the neighbours take $O(\text{adj}(v))$ time. // $\text{adj}(v)$ = adjacent nodes of v
- The sum of these over all vertices is $O(V^2 + E)$.
// Because $\sum \text{adj}(v) = E$ where E is the no of edges
- Can we improve this if we improve the running time for picking the **minimum distance()**? Yes, use **priority queue** to pick the **minimum**.

Using priority queue

foreach vertex w

$\text{distance}(w) = \text{INFINITY}$

$\text{distance}(s) = 0$

$\text{pq} = \text{new } \text{PriorityQueue}(V)$ // minimum heap

 //with all vertices and their distances (as keys)

while pq is not empty

$v = \text{pq.deleteMin}()$ // $O(\log V)$

foreach neighbour w of v

relax(v, w)

Since priority queue supports efficient minimum picking operation, we can use a priority queue here to improve the running time. Note that we no longer color vertices here. Yellow vertices in the previous pseudocode are now vertices that are in the priority queue.

Time Complexity - Initialization

foreach vertex w

$\text{distance}(w) = \text{INFINITY}$

$\text{distance}(s) = 0$

$\text{pq} = \text{new PriorityQueue}(V)$

 // with all vertices and their distances (as keys)

Initialization still takes $O(V)$

Time Complexity - Main loop

```
while pq is not empty  
    v = pq.deleteMin()  
    foreach neighbour w of v  
        relax(v,w)
```

We have to be more careful with the analysis of the main loop. We know that each deleteMin() takes $O(\log V)$ time. But relax(v,w) is no longer $O(1)$.

Note: Need to expand the relax(v,w) using priority queue

Time Complexity - Main loop


```
while pq is not empty
    v = pq.deleteMin()                // O(log V)
    foreach neighbour w of v          // adj(v)
        d = distance(v) + weight(v,w)
        if distance(w) > d then
            distance(w) = d
            pq.decreaseKey(w, d)      // O(log V)
            parent(w) = v
```

- If we expand the code for `relax()`, we will see that we **cannot** simply update `distance(v)`, since `distance(v)` is a key in the priority queue.
- Here, we use an operation called `decreaseKey()` that updates the key value of `distance(v)` in the priority queue.


Time Complexity - Main loop

- `decreaseKey()` can be done in $O(\log V)$ time. How?
- The time complexity for this version of Dijkstra's algorithm takes:
 - = $\text{sum } (O(\log V) + \text{adj}(v) * O(\log V))$ over all vertices
 - = $O(V \log V + E \log V) = O((V+E) \log V)$since the total number of adjacent nodes of all nodes is the total number of edges.

Dijkstra's algorithm




Dijkstra's algorithm



Vertex	Dist.
S	0

Step 1: Add source

Dijkstra's algorithm

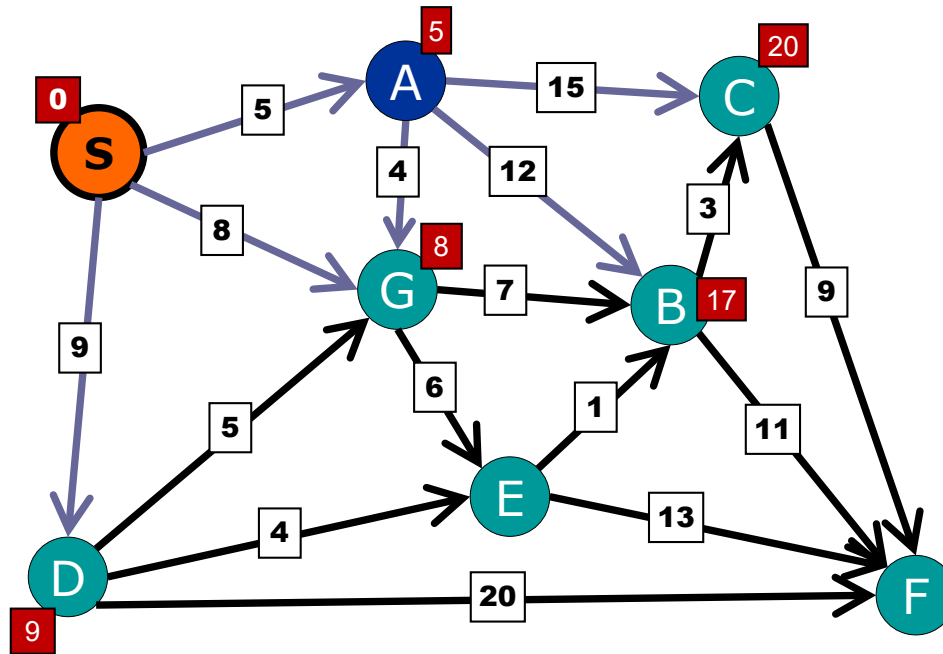


Vertex	Dist.
A	5
G	8
D	9

Step 1: Add source

Step 2: Remove S and relax.

Dijkstra's algorithm



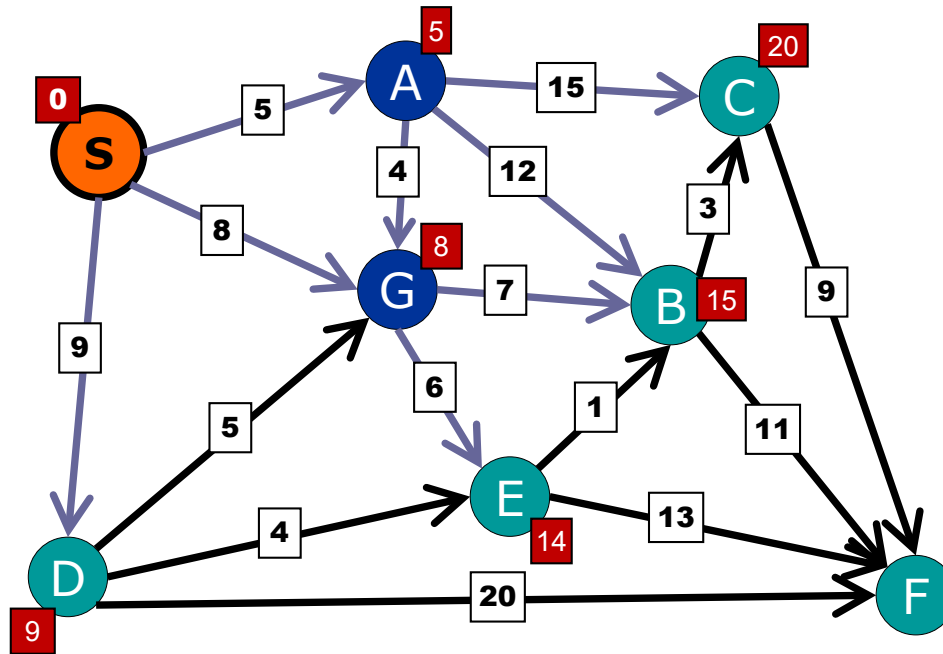
Vertex	Dist.
G	8
D	9
B	17
C	20

Step 1: Add source

Step 2: Remove S and relax.

Step 3: Remove A and relax.

Dijkstra's algorithm



Vertex	Dist.
D	9
E	14
B	15
C	20

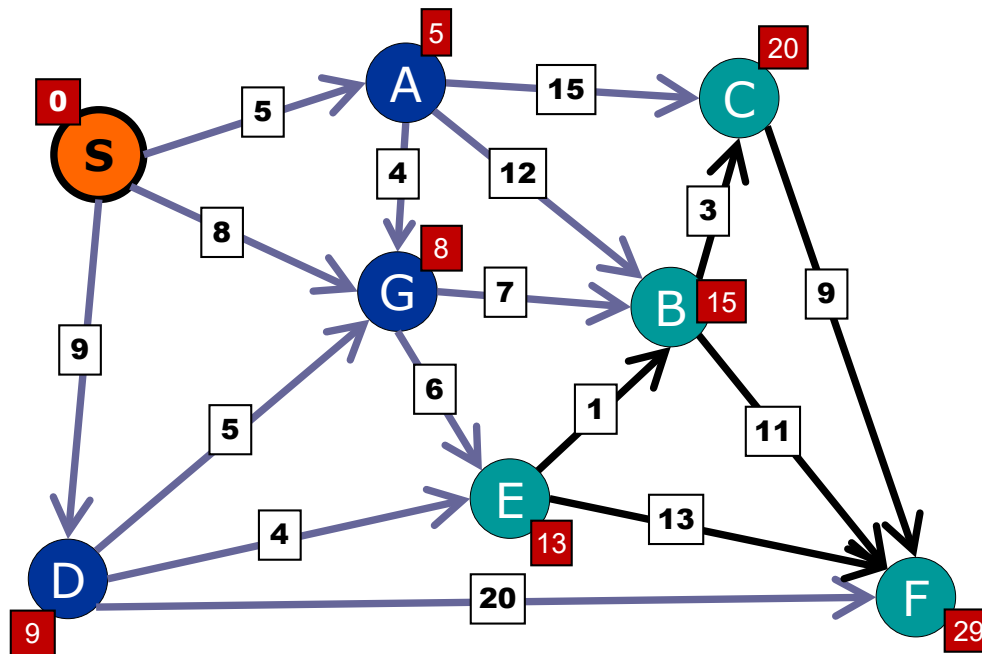
Step 1: Add source

Step 2: Remove S and relax.

Step 3: Remove A and relax.

Step 4: Remove G and relax.

Dijkstra's algorithm



Vertex	Dist.
E	13
B	15
C	20
F	29

Step 1: Add source

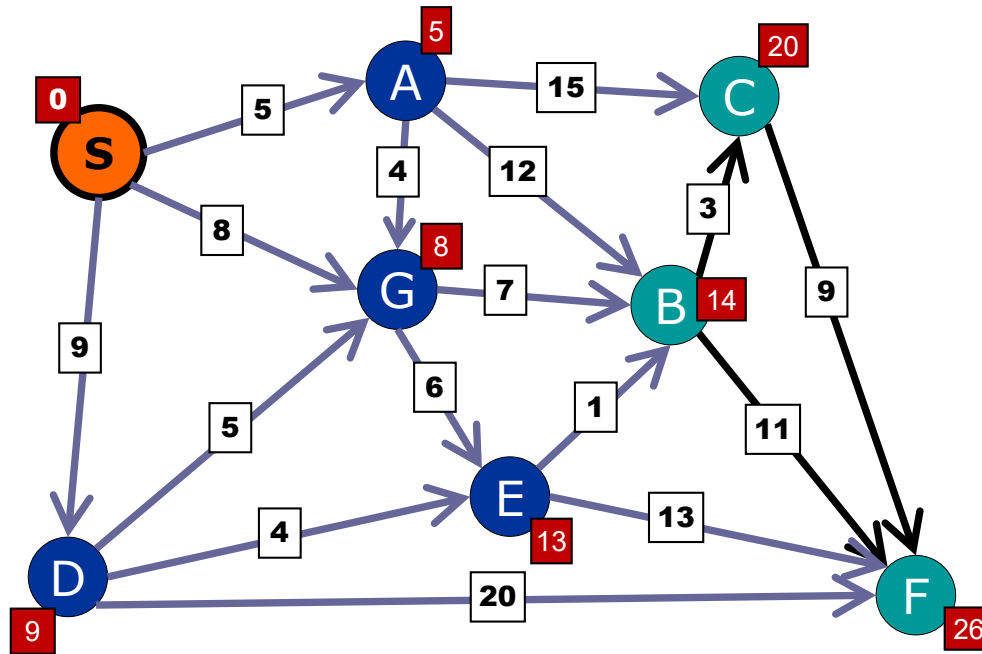
Step 2: Remove S and relax.

Step 3: Remove A and relax.

Step 4: Remove G and relax.

Step 5: Remove D and relax.

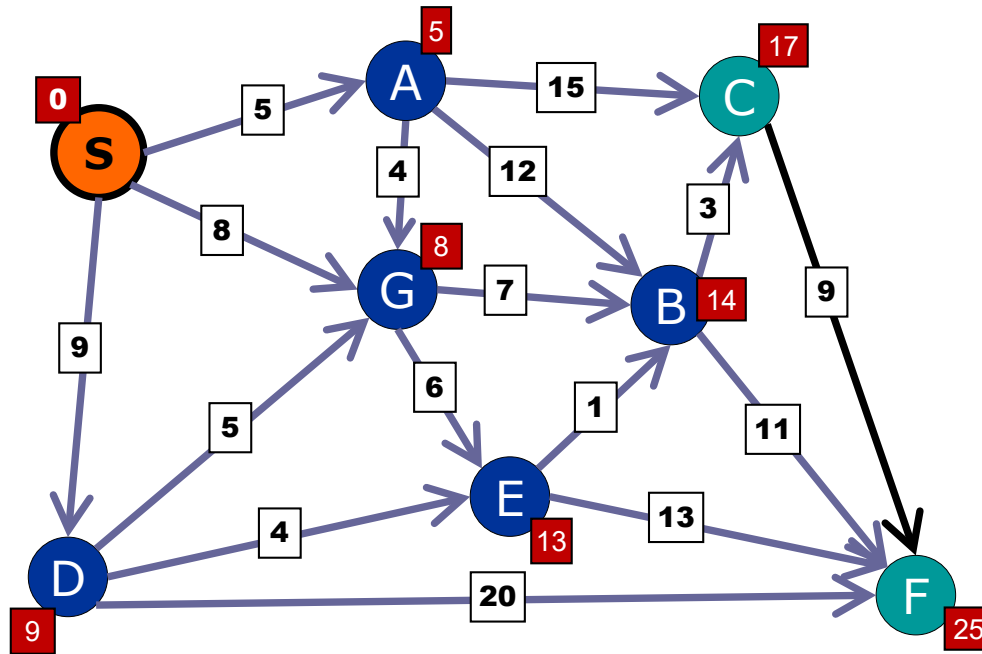
Dijkstra's algorithm



Vertex	Dist.
B	14
C	20
F	26

Step 6: Remove E and relax.

Dijkstra's algorithm

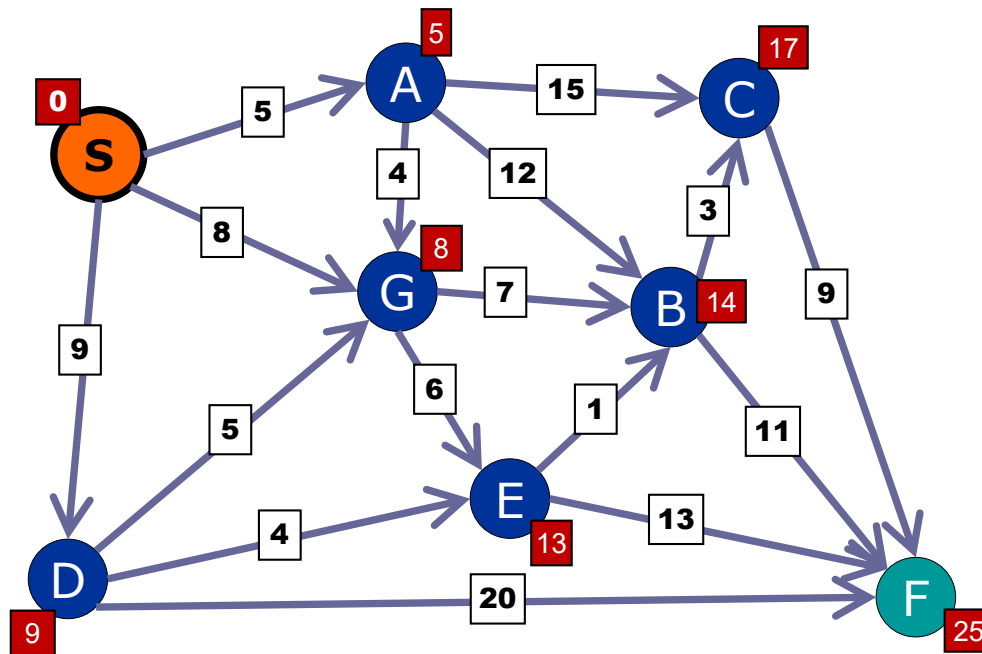


Vertex	Dist.
C	17
F	25

Step 6: Remove E and relax.

Step 7: Remove B and relax.

Dijkstra's algorithm



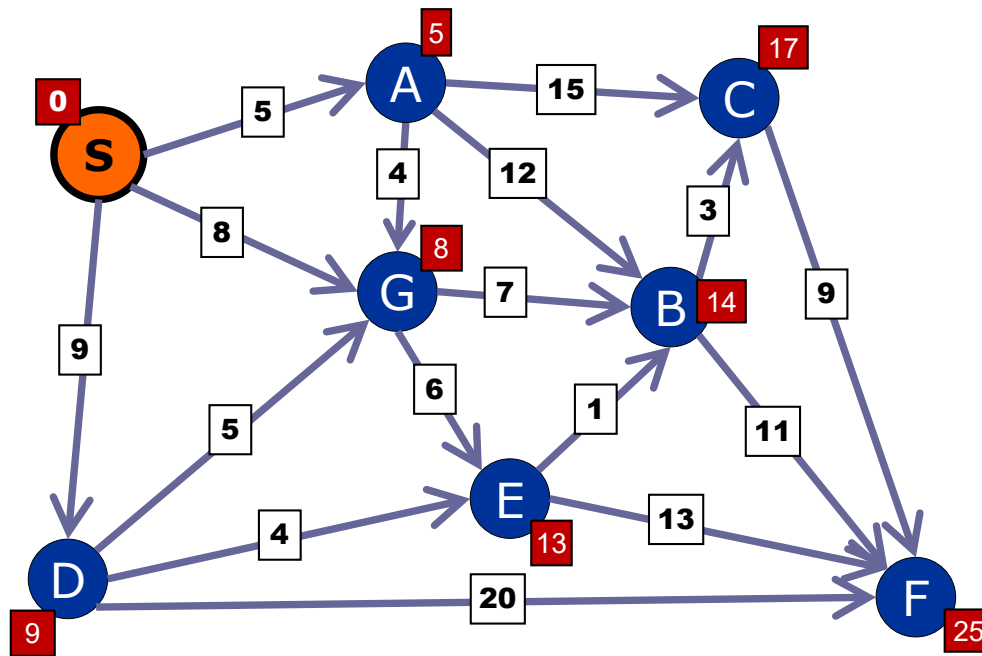
Vertex	Dist.
F	25

Step 6: Remove E and relax.

Step 7: Remove B and relax.

Step 8: Remove C and relax.

Dijkstra's algorithm



Vertex	Dist.

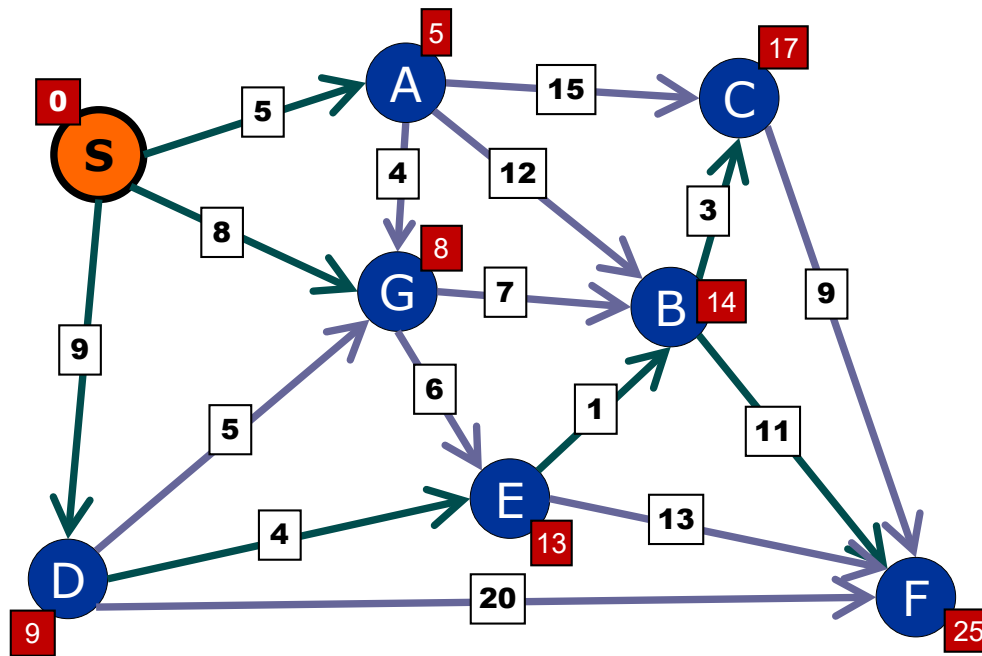
Step 6: Remove E and relax.

Step 7: Remove B and relax.

Step 8: Remove C and relax.

Step 9: Remove F and relax.

Dijkstra's algorithm



Vertex	Dist.

Step 6: Remove E and relax.


Step 7: Remove B and relax.

Step 8: Remove C and relax.

Step 9: Remove F and relax.

Done!

Dijkstra's algorithm




Vertex	Dist.

Step 10: Enjoy your
Shortest-Path Tree. 😊


Special Cases

Condition	Algorithm	Time Complexity
No Negative Weight Cycles	Bellman-Ford Algorithm	$O(VE)$
On Unweighted Graph (or equal weights)	BFS	$O(V + E)$
No Negative Weights	Dijkstra's Algorithm	
Negative Weights	Modified Dijkstra's Algorithm	
On Tree	BFS / DFS	$O(V)$
On DAG	Dynamic Programming	

Negative Weights



Negative Weights



Step 1: Remove A.
Relax A.
Mark A done.

...

Step 4: Remove B.
Relax B.
Mark B done.

Oops: We need to update A.

Dijkstra's algorithm does not work on graphs with negative weights


Modified Dijkstra's Algorithm

- Can be used for graphs with at least one negative weight edge
- Dijkstra's algorithm can also be implemented differently. The $O((V+E) \log V)$ **Modified Dijkstra's algorithm** can be used for directed weighted graphs that may have negative weight edges but no negative weight cycle.
- Such input graph appears in some practical cases, e.g. travelling using an **electric car** that has battery and our objective is to find a path from source vertex **s** to another vertex that minimizes overall **battery usage**. As usual, during acceleration (or driving on flat/uphill road), the electric car **uses** (positive) energy from the battery. However, during braking (or driving on downhill road), the electric car **recharges** (or use negative) energy to the battery. There is no negative weight cycle due to kinetic energy loss.

Modified Dijkstra's Algorithm

- The key idea is the modification done to C++ STL `priority_queue` to allow it to perform the required 'DecreaseKey' operation efficiently, i.e. in $O(\log V)$ time.
- The technique is called '*Lazy Update*' - leave the 'outdated/weaker/bigger-valued information' in the Min Priority Queue instead of deleting it straightaway. As the items are ordered from smaller values to bigger values in a Min PQ, we are guaranteeing ourselves that we will encounter the smallest/most-up-to-date item first before encountering the weaker/outdated item(s) later - which can be easily ignored.
- Refer to Visualgo for example

Negative Weights



Can we re-weight the edges with some constant (10)?


A. yeah!

B. Nope.. wouldn't work.

C.




Reweighting?



Path S-B-A: 1

Path S-A: 4

Reweighting?




Path S-B-A: 21

Path S-A: 14

The shortest path is no longer preserved!

Negative Weights



Can we re-weight the edges with some constant?

A. yeah!


B. Nope.. wouldn't work.




Special Cases

Condition	Algorithm	Time Complexity
No Negative Weight Cycles	Bellman-Ford Algorithm	$O(VE)$
On Unweighted Graph (or equal weights)	BFS	$O(V + E)$
No Negative Weights	Dijkstra's Algorithm	$O((V + E)\log V)$
Negative Weights	Modified Dijkstra's Algorithm	$O((V + E)\log V)$
On Tree	BFS / DFS	$O(V)$
On DAG	Dynamic Programming	

Directed Acyclic Graph (DAG)



Directed Acyclic Graph (DAG)




what relaxation order should we use for a DAG?

- A. Random.. any order is the same.
- B. BFS order
- C. Topological sort order

D.



Directed Acyclic Graph (DAG)



what relaxation order should we use for a DAG?

A. Random.. any order is the same.

B. BFS order


C. Topological sort order

D.




Directed Acyclic Graph (DAG)

1. Topological sort
2. Relax in order.




Directed Acyclic Graph (DAG)

1. Topological sort
2. Relax in order.




Directed Acyclic Graph (DAG)

1. Topological sort
2. Relax in order.




Directed Acyclic Graph (DAG)

1. Topological sort
2. Relax in order.




Directed Acyclic Graph (DAG)

1. Topological sort
2. Relax in order.




Directed Acyclic Graph (DAG)

1. Topological sort
2. Relax in order.




Directed Acyclic Graph (DAG)


1. Topological sort
2. Relax in order.



Why Topological Order?




Why Topological Order?



Fix S-D shortest path.


Why Topological Order?



Fix S-D shortest path.

S relaxed before A.


Why Topological Order?



Fix S-D shortest path.

S relaxed before A.
A relaxed before B.


Why Topological Order?



Fix S-D shortest path.

S relaxed before A.
A relaxed before B.
B relaxed before D.

Why Topological Order?



Fix S-D shortest path.

Path relaxed in-order,
so distance is correct
after relaxation.

Special Cases

Condition	Algorithm	Time Complexity
No Negative Weight Cycles	Bellman-Ford Algorithm	$O(VE)$
On Unweighted Graph (or equal weights)	BFS	$O(V + E)$
No Negative Weights	Dijkstra's Algorithm	$O((V + E)\log V)$
Negative Weights	Modified Dijkstra's Algorithm	$O((V + E)\log V)$
On Tree	BFS / DFS	$O(V)$
On DAG	Topological Sort (also called one-pass Bellman-Ford)	$O(V + E)$

Summary

- Looked at various SSSP algorithms for special graphs that can run faster
- Described each special graph and the algorithm used
- Described Dijkstra's algorithm and its modification
- Analyzed the computational complexity of Dijkstra's algorithm

*Acknowledgement: some slides courtesy of Dr Harold Soh